

AI Ethics Reading List

A short, credible reading list that moves from principles to practical risk management to current evidence.

How to read this list

- Do not try to read everything in one sitting.
- Start with one principles document, one practical framework, and one current trends report.
- For club discussion, the best question is not 'Do we like AI?' but 'What safeguards should exist for this use case?'

Suggested reading order

Priority	Reading	Why it matters	Best use
1	NIST AI Risk Management Framework	A broad official framework for thinking about trustworthy AI and managing risks across the AI lifecycle.	Big-picture structure
2	NIST AI RMF Playbook	Turns the framework into more actionable suggestions across govern, map, measure, and manage.	Practical follow-up
3	NIST Generative AI Profile	Extends the AI RMF specifically for generative AI and its newer risk patterns.	GenAI-specific discussion
4	OECD AI Principles	A compact international principles document updated in 2024 to address newer AI developments.	Principles and policy
5	UNESCO Recommendation on the Ethics of AI	Global human-rights-centered ethics guidance with policy action areas.	Values and safeguards

6	Stanford HAI 2025 AI Index Report	Strong annual overview of what is changing in AI performance, governance, responsible AI, economy, and public opinion.	Current evidence and trends
---	-----------------------------------	--	-----------------------------

What to watch for while reading

- Transparency: do people know when AI is being used, and what it can and cannot do?
- Human oversight: can a person override, review, or appeal an AI-supported decision?
- Bias and fairness: who gets harmed if the system is wrong more often for some groups?
- Privacy and data governance: where did the data come from and how is it being used?
- Safety and misuse: what happens if the system is used outside its intended purpose?
- Accountability: when something goes wrong, who is responsible for fixing it?

Good discussion questions for an AI club

- Which risks here are technical, and which are social or institutional?
- What kind of evidence would convince us an AI system is safe enough for a high-stakes use?
- Is the problem mainly model quality, weak governance, poor incentives, or all three?
- What should always require a human in the loop?
- What should users be told up front before they rely on the system?

Fast pairing guide

If you only read two things first

- Read NIST AI RMF for structure.
- Read Stanford HAI 2025 AI Index for current evidence.

Sources used

- NIST - AI Risk Management Framework
- NIST - AI RMF Playbook (updated 2025)
- NIST - AI RMF: Generative AI Profile (2024)
- OECD - AI Principles (updated in 2024)
- UNESCO - Recommendation on the Ethics of Artificial Intelligence (updated 2024)
- Stanford HAI - 2025 AI Index Report