

Prompting Cheat Sheet

A practical, model-agnostic guide for writing better prompts and getting more reliable outputs.

Core rule

- A better prompt does not mean a longer prompt.
- The goal is clearer instructions, clearer context, and clearer success criteria.
- Treat prompting as an iterative design task, not magic wording.

1. The basic prompt formula

Use this pattern

- Task: what the model should do
- Context: what the model needs to know
- Constraints: what must be true about the answer
- Format: how the answer should be structured
- Examples: optional, but useful if the target style is specific

2. High-value habits

Habit	Why it helps
Be clear and direct.	Official docs from Anthropic and OpenAI both emphasize clear, direct instructions instead of vague requests.
Give the model the goal, not just the topic.	'Explain transformers to Grade 10 students in 180 words' is much better than 'transformers'.
Specify the output format.	Ask for bullets, a table, a rubric, JSON, or a one-paragraph summary if that is what you actually need.
Add constraints early.	If there is a word limit, tone requirement, audience level, or allowed-source rule, say it explicitly.
Use examples when the target style is narrow.	Few-shot examples help when format and tone

	matter more than raw facts.
Iterate on weak outputs.	Google's prompt iteration guidance stresses changing and testing prompts in small steps.

3. Simple prompt templates

Explainer template:

Template

- Explain [topic] to [audience].
- Use [tone/style].
- Keep it to [length].
- Include [specific must-have elements].
- Do not include [specific unwanted elements].

Comparison template:

Template

- Compare [A] and [B] for [purpose].
- Use a table with columns for [criteria].
- End with a short recommendation for [audience/use case].

Research helper template:

Template

- Summarize the key points from these sources.
- Separate confirmed facts from inference.
- Flag any claim that still needs verification.

4. Prompting for reasoning models

- Do not assume you need 'think step by step' every time. OpenAI's current reasoning best-practices page says these models often work best with straightforward prompts.
- Keep the request crisp.
- Use delimiters, headings, or sections when the input has multiple parts.
- Start zero-shot first. Add examples only if needed.

- Be very specific about the end goal and success criteria.

5. Reliability checks

- Ask for sources when factual accuracy matters.
- Ask the model to separate known facts from uncertainty.
- Test your prompt on multiple cases, not just one lucky example.
- Use eval-style thinking: what would count as a bad answer, and can you detect it quickly?

6. A weak prompt vs a stronger prompt

Weak	Stronger
Tell me about neural networks.	Explain neural networks to Grade 10 students in under 180 words. Use one concrete example and avoid advanced math.
Make a slide deck on AI ethics.	Create a 5-slide outline for an AI ethics presentation aimed at high school students. Each slide should have a title and 3 bullets.
Summarize this article.	Summarize this article in 6 bullets. Then list 3 claims that should still be verified from the original source.

Sources used

- OpenAI API Docs - Prompting
- OpenAI API Docs - Reasoning best practices
- OpenAI API Docs - Evaluation best practices
- Anthropic Docs - Prompt engineering overview
- Google Cloud Vertex AI - Introduction to prompting
- Google Cloud Vertex AI - Prompt iteration strategies